

6.2a

July 13, 2021

1 Assignment 6.2a

Using section 5.2 in Deep Learning with Python as a guide, create a ConvNet model that classifies images CIFAR10 small images classification dataset. Do not use dropout or data-augmentation in this part. Save the model, predictions, metrics, and validation plots in the dsc650/assignments/assignment06/results directory. If you are using JupyterHub, you can include those plots in your Jupyter notebook.

```
[1]: from keras.datasets import cifar10
     from keras.utils import to_categorical
     import pandas as pd
```

```
[2]: (x_train, y_train), (x_test, y_test) = cifar10.load_data()
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>
170500096/170498071 [=====] - 18s 0us/step

```
[3]: x_train.shape, y_train.shape
```

```
[3]: ((50000, 32, 32, 3), (50000, 1))
```

```
[4]: x_test.shape, y_test.shape
```

```
[4]: ((10000, 32, 32, 3), (10000, 1))
```

```
[5]: # preprocess data
     x_train = x_train.astype('float32') / 255
     x_test = x_test.astype('float32') / 255
     y_train = to_categorical(y_train)
     y_test = to_categorical(y_test)

     # put 10,000 aside for validation
     x_val = x_train[-10000:]
     y_val = y_train[-10000:]
     x_train = x_train[:-10000]
     y_train = y_train[:-10000]
```

```
[6]: x_val.shape, y_val.shape
```

```
[6]: ((10000, 32, 32, 3), (10000, 10))
```

```
[7]: from keras import models
      from keras import layers
```

```
[8]: # instantiate the model
      model = models.Sequential()
      model.add(layers.Conv2D(32, (3,3), activation='relu', input_shape=(32, 32, 3)))
      model.add(layers.MaxPooling2D((2, 2)))
      model.add(layers.Conv2D(64, (3,3), activation='relu'))
      model.add(layers.MaxPooling2D((2,2)))
      model.add(layers.Conv2D(64, (3,3), activation='relu'))
      model.add(layers.MaxPooling2D((2,2)))
      model.add(layers.Flatten())
      model.add(layers.Dense(64, activation='relu'))
      model.add(layers.Dense(10, activation = 'softmax'))

      # view summary
      model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_1 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 6, 6, 64)	0
conv2d_2 (Conv2D)	(None, 4, 4, 64)	36928
max_pooling2d_2 (MaxPooling2D)	(None, 2, 2, 64)	0
flatten (Flatten)	(None, 256)	0
dense (Dense)	(None, 64)	16448
dense_1 (Dense)	(None, 10)	650
Total params: 73,418		
Trainable params: 73,418		
Non-trainable params: 0		

```
[9]: model.compile(optimizer='rmsprop',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])


history = model.fit(x_train, y_train, epochs=100,
                   validation_data = (x_val, y_val), verbose=0)
```

```
[10]: import matplotlib.pyplot as plt
```

```
[11]: train_loss = history.history['loss']
      val_loss = history.history['val_loss']

      epochs = range(1, len(history.history['loss']) + 1)

      plt.plot(epochs, train_loss, 'bo', label = 'Training Loss')
      plt.plot(epochs, val_loss, 'b', label = 'Validation Loss')
      plt.title('Training & Validation Losses')
      plt.xlabel('Epochs')
      plt.ylabel('Loss')
      plt.legend()
      plt.show()
      plt.savefig('results/6_2A_Loss.png')
```




<Figure size 432x288 with 0 Axes>

```
[12]: train_acc = history.history['accuracy']
      val_acc = history.history['val_accuracy']

      epochs = range(1, len(history.history['accuracy']) + 1)

      plt.plot(epochs, train_acc, 'bo', label = 'Training Accuracy')
      plt.plot(epochs, val_acc, 'b', label = 'Validation Accuracy')
      plt.title('Training & Validation Accuracy')
      plt.xlabel('Epochs')
      plt.ylabel('Accuracy')
      plt.legend()
      plt.show()
      plt.savefig('results/6_2A_Accuracy')
```



<Figure size 432x288 with 0 Axes>

```
[13]: # retrain model & evaluate
      (x_train, y_train), (x_test, y_test) = cifar10.load_data()

      # preprocess data
      x_train = x_train.astype('float32') / 255
      x_test = x_test.astype('float32') / 255
      y_train = to_categorical(y_train)
      y_test = to_categorical(y_test)
```

```

model.compile(optimizer='rmsprop',
              loss='categorical_crossentropy',
              metrics=['accuracy'])

history = model.fit(x_train, y_train, epochs=10)
results = model.evaluate(x_test, y_test)

```

```

Epoch 1/10
1563/1563 [=====] - 17s 10ms/step - loss: 0.7951 -
accuracy: 0.7852
Epoch 2/10
1563/1563 [=====] - 16s 10ms/step - loss: 0.6774 -
accuracy: 0.7976
Epoch 3/10
1563/1563 [=====] - 16s 10ms/step - loss: 0.6282 -
accuracy: 0.8066
Epoch 4/10
1563/1563 [=====] - 16s 10ms/step - loss: 0.6002 -
accuracy: 0.8117
Epoch 5/10
1563/1563 [=====] - 16s 10ms/step - loss: 0.5938 -
accuracy: 0.8113
Epoch 6/10
1563/1563 [=====] - 16s 10ms/step - loss: 0.5676 -
accuracy: 0.8212
Epoch 7/10
1563/1563 [=====] - 16s 10ms/step - loss: 0.5466 -
accuracy: 0.8238
Epoch 8/10
1563/1563 [=====] - 16s 10ms/step - loss: 0.5554 -
accuracy: 0.8224
Epoch 9/10
1563/1563 [=====] - 16s 10ms/step - loss: 0.5545 -
accuracy: 0.8162
Epoch 10/10
1563/1563 [=====] - 16s 10ms/step - loss: 0.5507 -
accuracy: 0.8233
313/313 [=====] - 1s 4ms/step - loss: 1.5735 -
accuracy: 0.6725

```

```
[14]: model.save('results/6_2A_model.h5')
```

```
[15]: prediction_results = model.predict(x_test)
```

```
[16]: prediction_results
```

```
[16]: array([[1.45420117e-05, 4.81691586e-06, 3.56239570e-06, ...,
            2.17049255e-05, 1.09994132e-02, 1.62153709e-04],
            [8.13679828e-04, 1.10762246e-01, 1.37421219e-14, ...,
            3.79440843e-14, 8.88385475e-01, 3.86396896e-05],
            [4.37420234e-03, 1.20730232e-02, 3.81997634e-05, ...,
            1.58175232e-03, 9.59185123e-01, 1.21063441e-02],
            ...,
            [3.28650163e-11, 5.68203995e-16, 3.56616511e-07, ...,
            1.60563843e-06, 3.80617282e-12, 7.23070701e-11],
            [1.07275538e-01, 1.57782082e-02, 5.50686538e-01, ...,
            2.40683369e-03, 2.12981395e-05, 1.66484058e-01],
            [0.00000000e+00, 0.00000000e+00, 2.00358356e-25, ...,
            1.00000000e+00, 0.00000000e+00, 0.00000000e+00]], dtype=float32)
```

```
[17]: with open('results/6_2A_metrics.txt', 'w') as f:
        f.write('Training Loss: {}'.format(str(history.history['loss'])))
        f.write('\nTraining Accuracy: {}'.format(str(history.history['accuracy'])))
        f.write('\nTest Loss: {}'.format(results[0]))
        f.write('\nTest Accuracy: {}'.format(results[1]))
```

```
[18]: preds = pd.DataFrame(prediction_results,
                           columns = ['0', '1', '2', '3', '4', '5', '6', '7', '8', '9'])

preds.to_csv('results/6_2A_predicitons.csv', index = False)
```

```
[ ]:
```