

## 6.2b

July 13, 2021

### 1 Assignment 6.2b

Using section 5.2 in Deep Learning with Python as a guide, create a ConvNet model that classifies images CIFAR10 small images classification dataset. This time includes dropout and data-augmentation. Save the model, predictions, metrics, and validation plots in the dsc650/assignments/assignment06/results directory. If you are using JupyterHub, you can include those plots in your Jupyter notebook.

```
[2]: from keras.datasets import cifar10
      from keras.utils import to_categorical
      from keras.preprocessing.image import ImageDataGenerator
      import pandas as pd
      import matplotlib.pyplot as plt
```

```
[3]: (x_train, y_train), (x_test, y_test) = cifar10.load_data()
```

```
[4]: x_train.shape, y_train.shape
```

```
[4]: ((50000, 32, 32, 3), (50000, 1))
```

```
[5]: x_test.shape, y_test.shape
```

```
[5]: ((10000, 32, 32, 3), (10000, 1))
```

```
[6]: # preprocess data
      x_train = x_train.astype('float32')
      x_test = x_test.astype('float32')
      y_train = to_categorical(y_train)
      y_test = to_categorical(y_test)

      # put 10,000 aside for validation
      x_val = x_train[-10000:]
      y_val = y_train[-10000:]
      x_train2 = x_train[:-10000]
      y_train2 = y_train[:-10000]
```

```
[7]: train_datagen = ImageDataGenerator(rescale=1./255,
                                         rotation_range=40,
```

```

        width_shift_range=0.2,
        height_shift_range=0.2,
        shear_range=0.2,
        zoom_range=0.2,
        horizontal_flip=True)

test_datagen = ImageDataGenerator(rescale=1./255)

train_generator = train_datagen.flow(x_train2, y_train2, batch_size=32)

validation_generator = train_datagen.flow(x_val, y_val, batch_size=32)

```

```

[8]: from keras import models
     from keras import layers

```

```

[9]: # instantiate the model
     # add dropout layer
     model = models.Sequential()
     model.add(layers.Conv2D(32, (3,3), activation='relu', input_shape=(32, 32, 3)))
     model.add(layers.MaxPooling2D((2, 2)))
     model.add(layers.Conv2D(64, (3,3), activation='relu'))
     model.add(layers.MaxPooling2D((2,2)))
     model.add(layers.Conv2D(64, (3,3), activation='relu'))
     model.add(layers.MaxPooling2D((2,2)))
     model.add(layers.Flatten())
     model.add(layers.Dropout(0.5))
     model.add(layers.Dense(64, activation='relu'))
     model.add(layers.Dense(10, activation = 'softmax'))

     # view summary
     model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_1 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 6, 6, 64)	0
conv2d_2 (Conv2D)	(None, 4, 4, 64)	36928
max_pooling2d_2 (MaxPooling2D)	(None, 2, 2, 64)	0

```

-----
flatten (Flatten)                (None, 256)                0
-----
dropout (Dropout)                (None, 256)                0
-----
dense (Dense)                    (None, 64)                 16448
-----
dense_1 (Dense)                  (None, 10)                 650
=====
Total params: 73,418
Trainable params: 73,418
Non-trainable params: 0
-----

```

```
[10]: from keras import optimizers
```

```
[11]: model.compile(optimizer=optimizers.RMSprop(lr=1e-4),
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])
```

```
[12]: history = model.fit_generator(train_generator,
                                   steps_per_epoch=len(x_train2) / 32,
                                   epochs = 30,
                                   validation_data=validation_generator,
                                   validation_steps=len(x_val) / 32)
```

```

/opt/conda/lib/python3.8/site-
packages/tensorflow/python/keras/engine/training.py:1844: UserWarning:
`Model.fit_generator` is deprecated and will be removed in a future version.
Please use `Model.fit`, which supports generators.
  warnings.warn("`Model.fit_generator` is deprecated and '

```

```

Epoch 1/30
1250/1250 [=====] - 44s 35ms/step - loss: 2.2275 -
accuracy: 0.1532 - val_loss: 1.9763 - val_accuracy: 0.2590
Epoch 2/30
1250/1250 [=====] - 43s 34ms/step - loss: 1.9872 -
accuracy: 0.2551 - val_loss: 1.9239 - val_accuracy: 0.2812
Epoch 3/30
1250/1250 [=====] - 43s 35ms/step - loss: 1.9184 -
accuracy: 0.2807 - val_loss: 1.8443 - val_accuracy: 0.3066
Epoch 4/30
1250/1250 [=====] - 43s 35ms/step - loss: 1.8564 -
accuracy: 0.3020 - val_loss: 1.8026 - val_accuracy: 0.3318
Epoch 5/30
1250/1250 [=====] - 43s 34ms/step - loss: 1.8197 -
accuracy: 0.3224 - val_loss: 1.7838 - val_accuracy: 0.3453
Epoch 6/30

```

1250/1250 [=====] - 43s 34ms/step - loss: 1.7992 - accuracy: 0.3302 - val\_loss: 1.7359 - val\_accuracy: 0.3638  
Epoch 7/30  
1250/1250 [=====] - 43s 34ms/step - loss: 1.7624 - accuracy: 0.3462 - val\_loss: 1.7389 - val\_accuracy: 0.3684  
Epoch 8/30  
1250/1250 [=====] - 43s 34ms/step - loss: 1.7322 - accuracy: 0.3588 - val\_loss: 1.6830 - val\_accuracy: 0.3871  
Epoch 9/30  
1250/1250 [=====] - 43s 35ms/step - loss: 1.7167 - accuracy: 0.3698 - val\_loss: 1.6856 - val\_accuracy: 0.3854  
Epoch 10/30  
1250/1250 [=====] - 43s 34ms/step - loss: 1.6997 - accuracy: 0.3772 - val\_loss: 1.6819 - val\_accuracy: 0.3895  
Epoch 11/30  
1250/1250 [=====] - 43s 34ms/step - loss: 1.6839 - accuracy: 0.3857 - val\_loss: 1.6256 - val\_accuracy: 0.4099  
Epoch 12/30  
1250/1250 [=====] - 43s 34ms/step - loss: 1.6568 - accuracy: 0.3973 - val\_loss: 1.6224 - val\_accuracy: 0.4113  
Epoch 13/30  
1250/1250 [=====] - 43s 34ms/step - loss: 1.6414 - accuracy: 0.4020 - val\_loss: 1.6037 - val\_accuracy: 0.4183  
Epoch 14/30  
1250/1250 [=====] - 42s 34ms/step - loss: 1.6364 - accuracy: 0.4049 - val\_loss: 1.5827 - val\_accuracy: 0.4258  
Epoch 15/30  
1250/1250 [=====] - 43s 35ms/step - loss: 1.6149 - accuracy: 0.4106 - val\_loss: 1.5742 - val\_accuracy: 0.4337  
Epoch 16/30  
1250/1250 [=====] - 43s 35ms/step - loss: 1.6177 - accuracy: 0.4142 - val\_loss: 1.5416 - val\_accuracy: 0.4440  
Epoch 17/30  
1250/1250 [=====] - 43s 35ms/step - loss: 1.6007 - accuracy: 0.4180 - val\_loss: 1.5630 - val\_accuracy: 0.4442  
Epoch 18/30  
1250/1250 [=====] - 43s 34ms/step - loss: 1.5854 - accuracy: 0.4289 - val\_loss: 1.5411 - val\_accuracy: 0.4532  
Epoch 19/30  
1250/1250 [=====] - 42s 34ms/step - loss: 1.5753 - accuracy: 0.4300 - val\_loss: 1.5057 - val\_accuracy: 0.4593  
Epoch 20/30  
1250/1250 [=====] - 43s 34ms/step - loss: 1.5696 - accuracy: 0.4346 - val\_loss: 1.5134 - val\_accuracy: 0.4594  
Epoch 21/30  
1250/1250 [=====] - 43s 34ms/step - loss: 1.5556 - accuracy: 0.4390 - val\_loss: 1.4859 - val\_accuracy: 0.4671  
Epoch 22/30

```

1250/1250 [=====] - 43s 34ms/step - loss: 1.5516 -
accuracy: 0.4496 - val_loss: 1.4945 - val_accuracy: 0.4653
Epoch 23/30
1250/1250 [=====] - 43s 34ms/step - loss: 1.5313 -
accuracy: 0.4507 - val_loss: 1.4835 - val_accuracy: 0.4711
Epoch 24/30
1250/1250 [=====] - 43s 34ms/step - loss: 1.5294 -
accuracy: 0.4515 - val_loss: 1.4549 - val_accuracy: 0.4774
Epoch 25/30
1250/1250 [=====] - 43s 34ms/step - loss: 1.5158 -
accuracy: 0.4554 - val_loss: 1.4437 - val_accuracy: 0.4837
Epoch 26/30
1250/1250 [=====] - 43s 34ms/step - loss: 1.5110 -
accuracy: 0.4578 - val_loss: 1.4439 - val_accuracy: 0.4864
Epoch 27/30
1250/1250 [=====] - 43s 34ms/step - loss: 1.5045 -
accuracy: 0.4573 - val_loss: 1.4495 - val_accuracy: 0.4853
Epoch 28/30
1250/1250 [=====] - 43s 34ms/step - loss: 1.4981 -
accuracy: 0.4616 - val_loss: 1.4305 - val_accuracy: 0.4903
Epoch 29/30
1250/1250 [=====] - 43s 34ms/step - loss: 1.4855 -
accuracy: 0.4697 - val_loss: 1.4197 - val_accuracy: 0.5021
Epoch 30/30
1250/1250 [=====] - 43s 35ms/step - loss: 1.4827 -
accuracy: 0.4669 - val_loss: 1.4237 - val_accuracy: 0.4965

```


```

[13]: train_loss = history.history['loss']
      val_loss = history.history['val_loss']

      epochs = range(1, len(history.history['loss']) + 1)

      plt.plot(epochs, train_loss, 'bo', label = 'Training Loss')
      plt.plot(epochs, val_loss, 'b', label = 'Validation Loss')
      plt.title('Training & Validation Losses')
      plt.xlabel('Epochs')
      plt.ylabel('Loss')
      plt.legend()
      plt.show()
      plt.savefig('results/6_2B_Loss.png')

```




<Figure size 432x288 with 0 Axes>

```
[14]: train_acc = history.history['accuracy']
      val_acc = history.history['val_accuracy']

      epochs = range(1, len(history.history['accuracy']) + 1)

      plt.plot(epochs, train_acc, 'bo', label = 'Training Accuracy')
      plt.plot(epochs, val_acc, 'b', label = 'Validation Accuracy')
      plt.title('Training & Validation Accuracy')
      plt.xlabel('Epochs')
      plt.ylabel('Accuracy')
      plt.legend()
      plt.show()
      plt.savefig('results/6_2B_Accuracy')
```



<Figure size 432x288 with 0 Axes>

```
[15]: # retrain model & evaluate
train_generator = train_datagen.flow(x_train, y_train, batch_size=32)

model.compile(optimizer=optimizers.RMSprop(lr=1e-4),
              loss='categorical_crossentropy',
              metrics=['accuracy'])

history = model.fit_generator(train_generator,
                             steps_per_epoch=len(x_train) / 32,
                             epochs = 16)

results = model.evaluate(x_test, y_test)
```

```
Epoch 1/16
1562/1562 [=====] - 45s 28ms/step - loss: 1.4893 -
accuracy: 0.4694
Epoch 2/16
1562/1562 [=====] - 44s 28ms/step - loss: 1.4692 -
accuracy: 0.4742
Epoch 3/16
1562/1562 [=====] - 44s 28ms/step - loss: 1.4655 -
accuracy: 0.4750
```

```

Epoch 4/16
1562/1562 [=====] - 44s 28ms/step - loss: 1.4514 -
accuracy: 0.4830
Epoch 5/16
1562/1562 [=====] - 44s 28ms/step - loss: 1.4534 -
accuracy: 0.4832
Epoch 6/16
1562/1562 [=====] - 44s 28ms/step - loss: 1.4350 -
accuracy: 0.4857
Epoch 7/16
1562/1562 [=====] - 44s 28ms/step - loss: 1.4365 -
accuracy: 0.4890
Epoch 8/16
1562/1562 [=====] - 44s 28ms/step - loss: 1.4232 -
accuracy: 0.4934
Epoch 9/16
1562/1562 [=====] - 44s 28ms/step - loss: 1.4239 -
accuracy: 0.4897
Epoch 10/16
1562/1562 [=====] - 44s 28ms/step - loss: 1.4199 -
accuracy: 0.4930
Epoch 11/16
1562/1562 [=====] - 44s 28ms/step - loss: 1.4094 -
accuracy: 0.4991
Epoch 12/16
1562/1562 [=====] - 44s 28ms/step - loss: 1.3971 -
accuracy: 0.5016
Epoch 13/16
1562/1562 [=====] - 44s 28ms/step - loss: 1.3961 -
accuracy: 0.5012
Epoch 14/16
1562/1562 [=====] - 44s 28ms/step - loss: 1.3908 -
accuracy: 0.5060
Epoch 15/16
1562/1562 [=====] - 44s 28ms/step - loss: 1.3894 -
accuracy: 0.5045
Epoch 16/16
1562/1562 [=====] - 44s 28ms/step - loss: 1.3856 -
accuracy: 0.5065
313/313 [=====] - 1s 4ms/step - loss: 239.9552 -
accuracy: 0.3348

```

```
[16]: model.save('results/6_2B_model.h5')
```

```
[17]: prediction_results = model.predict(x_test)
```

```
[18]: prediction_results
```



```
[18]: array([[0.0000000e+00, 0.0000000e+00, 0.0000000e+00, ..., 0.0000000e+00,
            1.0000000e+00, 0.0000000e+00],
            [0.0000000e+00, 1.0000000e+00, 0.0000000e+00, ..., 0.0000000e+00,
            0.0000000e+00, 0.0000000e+00],
            [5.8108697e-17, 1.0000000e+00, 0.0000000e+00, ..., 0.0000000e+00,
            0.0000000e+00, 0.0000000e+00],
            ...,
            [0.0000000e+00, 0.0000000e+00, 0.0000000e+00, ..., 1.1507828e-14,
            0.0000000e+00, 0.0000000e+00],
            [0.0000000e+00, 1.0000000e+00, 0.0000000e+00, ..., 0.0000000e+00,
            0.0000000e+00, 0.0000000e+00],
            [0.0000000e+00, 0.0000000e+00, 0.0000000e+00, ..., 1.0000000e+00,
            0.0000000e+00, 0.0000000e+00]], dtype=float32)
```

```
[19]: with open('results/6_2B_metrics.txt', 'w') as f:
        f.write('Training Loss: {}'.format(str(history.history['loss'])))
        f.write('\nTraining Accuracy: {}'.format(str(history.history['accuracy'])))
        f.write('\nTest Loss: {}'.format(results[0]))
        f.write('\nTest Accuracy: {}'.format(results[1]))
```

```
[20]: preds = pd.DataFrame(prediction_results,
                           columns = ['0', '1', '2', '3', '4', '5', '6', '7', '8', '9'])

preds.to_csv('results/6_2B_predicitons.csv', index = False)
```

```
[ ]:
```